




Intelligent Clinical Early Warning System

Can AI Save Lives?

Name: Muhammad Sohaib

Roll No: 22F-BSAI-40

Department: BS-Artificial Intelligence

Submitted to: Engr. Hamza Farooqui

Section 1: The Clinical Problem

1.1 Why Patient Deterioration Differs from Standard Classification

Predicting patient deterioration in an ICU is fundamentally different from a conventional classification task. Three critical challenges define this problem domain:

- Extreme class imbalance: In the PhysioNet Sepsis 2019 dataset, only 1.80% of patient-hours carry a positive SepsisLabel. A naïve classifier that always predicts 'No Sepsis' achieves 98.2% accuracy while saving zero lives. Recall (Sensitivity) must therefore be the primary metric.
- Temporal structure and evolving risk: A patient's physiology is a narrative, not a snapshot. Static models cannot capture the progressive cascade (tachycardia → rising lactate → hypotension → organ failure) characterising septic shock. This motivated escalation from static DNNs (Generation 1) to temporal sequence models (Generation 2).
- Noisy and delayed labels: Ground-truth SepsisLabel is derived retrospectively, meaning the model must predict risk before the label formally fires — a fundamentally harder problem.
- Catastrophic missing data: Clinically important laboratory values are missing in over 95% of rows (e.g., EtCO2 at 96%, Fibrinogen at 99%), introducing systematic bias if not handled carefully.

Metric	Standard Task	Clinical Deterioration
Label quality	Clean, precise	Delayed, noisy, subjective
Class balance	Often balanced	1.8% positive (extreme imbalance)
Temporal structure	None required	Hours of sequential vitals needed
Cost of False Negative	Low	Patient mortality — unacceptable

Table 1: Clinical deterioration prediction vs. standard classification

1.2 Why Recall Matters More Than Accuracy

In a life-critical AI system, a missed sepsis case (False Negative) can cost a patient's life, while a false alarm (False Positive) costs only nursing time for an unnecessary assessment. This asymmetry is severe and drives all design decisions.

Example — Patient Maria Chen, 68: At hour 18 post-surgery, subtle vital-sign shifts begin. A high-accuracy but low-recall model (Recall = 45%) fails to alert; Maria goes into septic shock at hour 22, increasing her 30-day mortality by an estimated 7% per hour of untreated sepsis. A recall-optimised model (Recall = 69–91%) fires an alert at the correct time, enabling antibiotics 4 hours earlier with dramatically improved outcomes.

Formally: $\text{Recall} = \text{TP} / (\text{TP} + \text{FN})$. The project uses a decision threshold of 0.35 (rather than the default 0.5), deliberately accepting more false alarms to maximise sensitivity. Precision and Accuracy are secondary constraints, not primary objectives.

Section 2: Baseline Design Decisions

2.1 Vanishing Gradients, ReLU, and Batch Normalisation

Deep networks trained with backpropagation suffer from vanishing gradients: the product of Jacobian matrices across many layers shrinks exponentially, starving shallow layers of learning signal. The project addressed this with two measures:

- ReLU activations: $f(x) = \max(0, x)$ has a gradient of exactly 1 for all positive inputs — neither growing nor shrinking — directly preventing vanishing gradients in the active regime.

- Batch Normalisation: Applied after each Dense layer, BatchNorm normalises pre-activation values to zero mean and unit variance, preventing extreme saturation and acting as an implicit regulariser.

2.2 Optimizer and Loss Function

The same DNN was trained with Adam and SGD to compare performance:

Attribute	Adam	SGD + Momentum	Winner
Accuracy	0.7037	0.6653	Adam
Recall	0.6512	0.5965	Adam
F1-Score	0.6873	0.6405	Adam
Convergence	Smooth, stable	Oscillatory, early-stopped ep 8	Adam

Table 2: Adam vs SGD optimizer comparison (DNN Baseline, test set)

Adam was selected for all subsequent architectures due to its faster convergence, robustness to hyperparameter choice, and superior Recall. The loss function used was Binary Cross-Entropy (BCE); in PyTorch implementations, a `pos_weight` of 51.3 scaled the cost for positive (sepsis) examples, encoding the clinical priority of minimising false negatives directly into the mathematical objective.

Section 3: Modelling Patient Timelines

3.1 LSTM — Solving Temporal Vanishing Gradients

Vanilla RNNs suffer from vanishing gradients in the time domain: the gradient product across T timesteps attenuates exponentially, making dependencies beyond 5–7 steps unlearnable. In a 12-hour ICU window, this is catastrophic — early vital-sign trends that predict later sepsis onset become invisible to the model.

3.2 GRU — Preferred Architecture for Real-Time Deployment

The Gated Recurrent Unit (GRU) merges LSTM's forget and input gates into a single update gate, reducing parameters by 24% (162,945 vs 214,401) with marginally faster training. In experiments, GRU outperformed LSTM on the clinical task:

Architecture	Recall (Test)	AUC	Train Time	Deployable?
LSTM	0.6906	0.705	2.5 min	Yes
GRU	0.9092	0.733	2.4 min	YES — recommended
BiLSTM	0.8952	0.726	4.2 min	NO — future data leak

Table 3: Temporal architectures performance (threshold = 0.35)

The GRU's superior Recall is attributed to its fewer parameters being less prone to overfitting on the severely imbalanced dataset (1.8% positive rate). Its lighter footprint makes it ideal for resource-constrained bedside hardware.

3.3 Why BiLSTM Is Excluded from Deployment

A Bidirectional LSTM processes both past and future context simultaneously. For offline research, this improves performance — but for a live Early Warning System, it constitutes a fundamental data leak: the

model at hour 8 implicitly 'knows' what happens at hour 12, making the alert clinically meaningless. Only causal architectures (LSTM, GRU) are valid for real-time deployment.

Section 4: The Transformer and Clinical Language

4.1 ClinicalBERT and Transfer Learning

Vital-sign numbers cannot distinguish three patients with identical readings of HR = 112, O2Sat = 94%, RR = 22 — one post-operative and stable, one in septic shock, one with controlled atrial fibrillation. Unstructured clinical notes add contextual framing, qualitative symptoms, and treatment intent that numbers alone cannot capture.

ClinicalBERT (emilyalsentzer/Bio_ClinicalBERT), pre-trained on PubMed abstracts and MIMIC-III notes, was fine-tuned on 50,000 clinical notes. Transfer learning is essential: training 108 million parameters from scratch on a small dataset would fail to acquire clinical terminology, negation understanding ('no fever' ≠ 'fever'), or complex temporal phrasing.

4.2 Self-Attention and Fine-Tuning Strategy

The Transformer's self-attention mechanism allows every token to directly attend to every other token: $\text{Attention}(Q, K, V) = \text{softmax}(QK^T / \sqrt{d_k}) \cdot V$. This means 'sepsis' in sentence 3 can directly connect to 'lactate elevated' in sentence 7, regardless of intervening words — a key advantage over LSTMs for long clinical notes.

Strategy	Trainable Params	Wall-Clock	Val Accuracy	F1-Score
Frozen Base	197,378	52.1 s	0.6266	0.5942
Full Fine-Tune	108,507,650	97.6 s	0.6449	0.6169

Table 4: ClinicalBERT fine-tuning strategy comparison (3,000-note sample)

Full fine-tuning achieved 1.83 points better validation accuracy at 1.87× the training time. It is preferred when a GPU server is available and dataset size is sufficient to update 108M parameters without overfitting.

Section 5: Deployment Verdict

5.1 Recommended Production Architecture

Model	Recall	AUC	Latency	ICU Deployment?
DNN Baseline	0.651	N/A	<1 ms	No — no temporal context
LSTM	0.691	0.705	~5 ms	Yes — causal, good recall
GRU	0.909	0.733	~4 ms	YES — best recall, lightest
BiLSTM	0.895	0.726	~8 ms	NO — future data leak
ClinicalBERT (Frozen)	0.604	N/A	~50 ms	Limited — slow
ClinicalBERT (Full FT)	0.618	N/A	~50 ms	YES (notes) — best text model

Table 5: Comprehensive model comparison across all six architectures

The recommended production system is a two-component fusion architecture:

- Primary Vital-Signs Monitor — GRU: Highest Recall (90.9%) among causal models, AUC 0.733, 4 ms inference latency, deployable on resource-constrained bedside hardware without a dedicated GPU.
- Clinical Notes Risk Scorer — ClinicalBERT (Full Fine-Tune): Best text-based model with F1 of 0.617 and 50 ms latency per note (acceptable given notes are generated episodically every 2–4 hours).
- Fusion: $\text{Risk_Score_final} = \alpha \times \text{GRU_score} + (1-\alpha) \times \text{ClinicalBERT_score}$, with $\alpha \approx 0.6$. An alert fires when the fused score exceeds a calibrated threshold, reviewed annually against confirmed outcomes.

5.2 Ethical Considerations

Three critical ethical concerns must be addressed before any clinical deployment:

- Algorithmic Bias: Models trained on historical data inherit documented healthcare disparities — e.g., undertreated pain in Black and Hispanic patients, cardiac symptoms misattributed to anxiety in women. Stratified performance evaluation across demographic subgroups is mandatory before deployment.
- Patient Privacy: Clinical notes are among the most sensitive personal data available. All patient data must remain within the hospital's secure perimeter; no cloud inference is permissible without explicit consent and HIPAA/GDPR-compliant data processing agreements.
- Human-in-the-Loop: Every alert must be reviewed and acted upon by a licensed clinician. The model is a second opinion, not a decision-maker. Alert fatigue (from excessive false positives) and over-reliance are both risks requiring ongoing calibration and mandatory documentation of alert responses.

5.3 Conclusions

This project presented a three-generation pipeline for intelligent clinical early warning:

- Generation 1 (DNN): Static vital-sign modelling established core design choices (Adam, ReLU, BatchNorm) with 70.4% accuracy but limited temporal insight.
- Generation 2 (GRU/LSTM): Sequential modelling dramatically improved Recall — GRU reached 90.9% vs the DNN's 65.1%. Causal architectures validated as the only appropriate choice for real-time deployment.
- Generation 3 (ClinicalBERT): Unstructured clinical notes carry risk signal inaccessible to vital-sign models alone. Transfer learning from pre-trained clinical language models is essential.